

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sb
import matplotlib.pyplot as plt
import scipy as spy
import keras
from sklearn.metrics import accuracy_score, recall_score, confusion_matrix
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, Dropout, Input
from keras.optimizers import Adam, RMSprop
df = pd.read_csv('winequality-white.csv')
```

2021-08-11 12:11:24.330688: W tensorflow/stream\_executor/platform/default/dso\_loader.cc:64] Could not load dynamic library 'libcudart.so.11.0'; dlerror: libcudart.so.11.0: cannot open shared object file: No such file or directory  
 2021-08-11 12:11:24.330706: I tensorflow/stream\_executor/cuda/cudart\_stub.cc:29] Ignore above cudart dlerror if you do not have a GPU set up on your machine.

Olhando abaixo, não temos nenhum valor N/A, então não precisamos tratar isso.

```
In [2]: df.isna().sum()
```

```
Out[2]: fixed acidity      0
volatile acidity    0
citric acid         0
residual sugar      0
chlorides           0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

```
In [3]: print(len(df))
df
```

4898

```
Out[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.00100	3.00	0.45	8.8	6
1	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.99400	3.30	0.49	9.5	6
2	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.99510	3.26	0.44	10.1	6
3	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
4	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
...	...	...	...	...	...	...	...	...	...	...	...	...
4893	6.2	0.21	0.29	1.6	0.039	24.0	92.0	0.99114	3.27	0.50	11.2	6
4894	6.6	0.32	0.36	8.0	0.047	57.0	168.0	0.99490	3.15	0.46	9.6	5
4895	6.5	0.24	0.19	1.2	0.041	30.0	111.0	0.99254	2.99	0.46	9.4	6
4896	5.5	0.29	0.30	1.1	0.022	20.0	110.0	0.98869	3.34	0.38	12.8	7
4897	6.0	0.21	0.38	0.8	0.020	22.0	98.0	0.98941	3.26	0.32	11.8	6

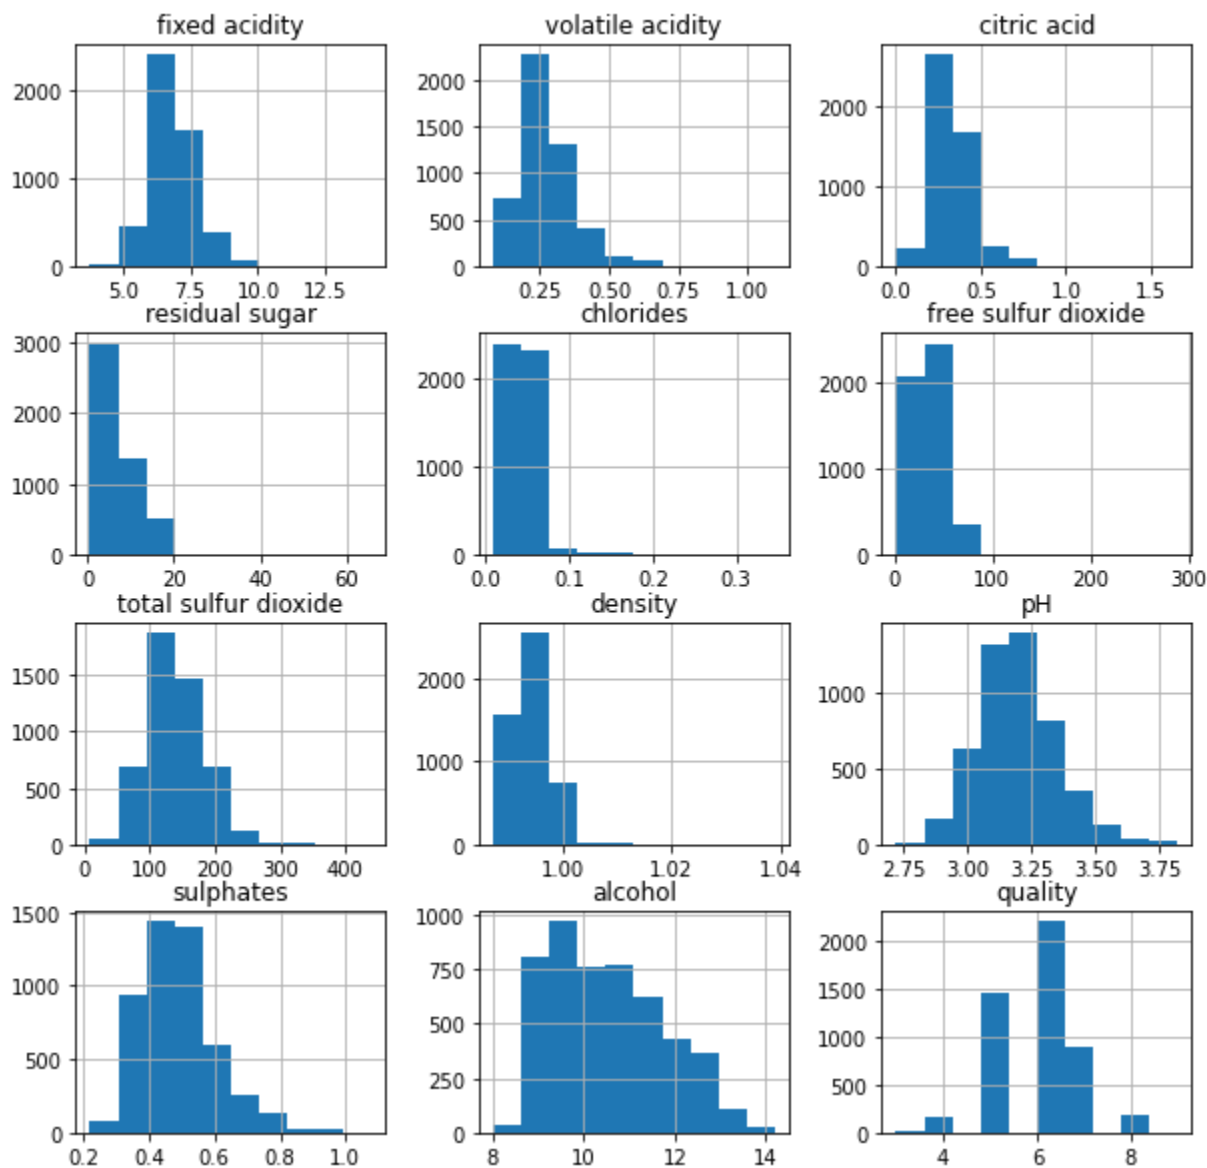
4898 rows × 12 columns

```
In [4]: df['quality'].value_counts()
```

```
Out[4]: 6      2198
        5      1457
        7       880
        8       175
        4       163
        3        20
        9         5
        Name: quality, dtype: int64
```

```
In [5]: df.hist(figsize = (10, 10))
```

```
Out[5]: array([[<AxesSubplot:title={'center':'fixed acidity'}>,
               <AxesSubplot:title={'center':'volatile acidity'}>,
               <AxesSubplot:title={'center':'citric acid'}>],
              [<AxesSubplot:title={'center':'residual sugar'}>,
               <AxesSubplot:title={'center':'chlorides'}>,
               <AxesSubplot:title={'center':'free sulfur dioxide'}>],
              [<AxesSubplot:title={'center':'total sulfur dioxide'}>,
               <AxesSubplot:title={'center':'density'}>,
               <AxesSubplot:title={'center':'pH'}>],
              [<AxesSubplot:title={'center':'sulphates'}>,
               <AxesSubplot:title={'center':'alcohol'}>,
               <AxesSubplot:title={'center':'quality'}>]], dtype=object)
```



Podemos ver acima que temos um desbalanceamento na quantidade de amostras. Temos muitas regulares e poucas ruins e ótimas. Vamos usar uma técnica chamada SMOTE que consiste em fazer o oversampling das amostras minoritárias, deixando assim o dataset balanceado. Essa técnica foi descrita no artigo.

```
In [6]: X=df.drop(columns=['quality'])
        y=df['quality']
```

```
In [7]: from imblearn.over_sampling import SMOTE
        oversample = SMOTE(k_neighbors=4)
        X, y = oversample.fit_resample(X, y)
```

```
In [8]: print(y.dtypes)
        print(y.count())
        y.value_counts()
```

```
Out[8]: int64
        15386
        6    2198
        5    2198
        7    2198
        8    2198
        4    2198
        3    2198
```


9 2198  
Name: quality, dtype: int64

Pronto, agora temos todas as amostras em quantias iguais.

mapa de correlação:

```
In [9]: corr=df.corr()  
plt.figure(figsize=(20,10))  
sb.heatmap(corr,annot=True, cmap='coolwarm')
```

Out[9]: <AxesSubplot:>



Também quero codificar essas qualidades numericamente:

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.20,random_state=21)  
print('Formato do dataset de treinamento Xs:{}'.format(X_train.shape))  
print('Formato do dataset de teste Xs:{}'.format(X_test.shape))  
print('Formato do dataset de treino y:{}'.format(y_train.shape))  
print('Formato do dataset de test y:{}'.format(y_test.shape))  
y_train
```

```
Formato do dataset de treinamento Xs:(12308, 11)  
Formato do dataset de teste Xs:(3078, 11)  
Formato do dataset de treino y:(12308,)  
Formato do dataset de test y:(3078,)
```

```
Out[10]: 12509      8  
14261      9  
13650      9  
7263       4  
2851       5  
..  
48         6  
8964       4  
5944       3  
5327       3
```

```
15305      9
Name: quality, Length: 12308, dtype: int64
```

Escalando o dataset

```
In [11]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train=sc.fit_transform(X_train)
X_test=sc.fit_transform(X_test)
X_train
```

```
Out[11]: array([[ -0.22741558, -1.0639914 , -0.38458136, ...,  0.00307363,
        -0.12161052,  0.06661161],
       [ -0.16436717,  0.57646457,  0.01817661, ...,  0.50304968,
        -1.14134359,  1.42526059],
       [ -0.13978815,  0.48381    ,  0.16594051, ...,  0.55187255,
        -1.11749413,  1.44325812],
       ...,
       [  1.04942297,  0.19851959,  0.8274273 , ...,  0.16877131,
        -1.44675857,  1.13961914],
       [  0.87318938, -0.4188991 , -1.07750024, ..., -0.04315277,
        0.59747948, -0.92846942],
       [  0.16432647, -0.45375517,  1.04106329, ...,  0.89852317,
        -0.39543528,  1.47342487]])
```

```
In [12]: from sklearn.ensemble import RandomForestClassifier
forest=RandomForestClassifier(n_estimators=10, criterion = 'entropy', random_state=0)
forest.fit(X_train,y_train)
```

```
Out[12]: RandomForestClassifier(criterion='entropy', n_estimators=10, random_state=0)
```

Outros testes com outro método

## avaliação do resultado:

```
In [15]: print('Acurácia da floresta: ',forest.score(X_test,y_test))
```

Acurácia da floresta: 0.8515269655620533

```
In [ ]: y_pred = forest.predict(X_test)
```

```
In [ ]: #R2 score
from sklearn.metrics import r2_score
r2_score(y_test, y_pred)
```